

References

- Abraham, W. C., & Robins, A. (2005). Memory retention - the synaptic stability versus plasticity dilemma. *Trends in Neurosciences*, 28(2), 73–78. <https://doi.org/10.1016/j.tins.2004.12.003>
- Agrawal, L. A., Tan, S., Soylu, D., Ziemis, N., Khare, R., Opsahl-Ong, K., Singhvi, A., Shandilya, H., Ryan, M. J., Jiang, M., Potts, C., Sen, K., Dimakis, A. G., Stoica, I., Klein, D., Zaharia, M., & Khattab, O. (2024). *GEPA: Reflective prompt evolution can outperform reinforcement learning*. arXiv. <https://doi.org/10.48550/arXiv.2507.19457>
- Aguiar, L., & Waldfogel, J. (2021). Platforms, power, and promotion: Evidence from spotify playlists. *Journal of Industrial Economics*, 69(3), 653–691. <https://doi.org/10.1111/joie.12263>
- Albu-Schäffer, A., Eiberger, O., Grebenstein, M., Haddadin, S., Ott, C., Wimböck, T., Wolf, S., & Hirzinger, G. (2008). Soft robotics. *IEEE Robotics & Automation Magazine*, 15(3), 20–30. <https://doi.org/10.1109/MRA.2008.927979>
- Alhakami, M., Ashley, D. R., Dunham, J., Dai, Y., Faccio, F., Feron, E., & Schmidhuber, J. (2025). Towards an extremely robust baby robot with rich interaction ability for advanced machine learning algorithms. *Proceedings of the 2025 IEEE/RSJ International Conference on Intelligent Robots and Systems*, 16725–16730. <https://doi.org/10.1109/IROS60139.2025.11246188>
- Amaral, J. N., Borin, E., Ashley, D. R., Benedicto, C., Colp, E., Hoffmam, J. H. S., Karpoff, M., Ochoa, E., Redshaw, M., & Rodrigues, R. E. (2018). The alberta workloads for the SPEC CPU 2017 benchmark suite. *Proceedings of the 2018 IEEE International Symposium on Performance Analysis of Systems and Software*, 159–168. <https://doi.org/10.1109/ISPASS.2018.00029>
- Andric, A., & Haus, G. (2006). Automatic playlist generation based on tracking user’s listening habits. *Multimedia Tools and Applications*, 29(2), 127–151. <https://doi.org/10.1007/S11042-006-0003-9>
- Andriushchenko, M., Croce, F., & Flammarion, N. (2024). *Jailbreaking leading safety-aligned LLMs with simple adaptive attacks*. arXiv. <https://doi.org/10.48550/arXiv.2404.02151>
Showed that different previously developed attack strategies work well on jailbreaking different LLMs and that trying them all gives a 100% success rate.
- Andrychowicz, M., Baker, B., Chociej, M., Józefowicz, R., McGrew, B., Pachocki, J., Petron, A., Plappert, M., Powell, G., Ray, A., Schneider, J., Sidor, S., Tobin, J., Welinder, P., Weng, L., & Zaremba, W. (2020). Learning dexterous in-hand manipulation. *The International Journal of Robotics Research*, 39(1). <https://doi.org/10.1177/0278364919887447>
- Andrychowicz, M., Crow, D., Ray, A., Schneider, J., Fong, R., Welinder, P., McGrew, B., Tobin, J., Abbeel, P., & Zaremba, W. (2017). Hindsight experience replay. *Advances in Neural Information Processing Systems*, 30, 5048–5058. https://proceedings.neurips.cc/paper_files/paper/2017/file/453fadbd8a1a3af50a9df4df899537b5-Paper.pdf
- Argaw, D., Soldan, M., Pardo, A., Zhao, C., Caba, F., Chung, J. S., & Ghanem, B. (2024). Towards automated trailer generation. *Proceedings of the 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition*.
- Arulkumaran, K., Ashley, D. R., Schmidhuber, J., & Srivastava, R. K. (2022). *All you need is supervised learning: From imitation learning to meta-RL with upside down RL*. arXiv. <https://doi.org/10.48550/arXiv.2202.11960>
- Ashley, D. R. (2020). *Understanding Forgetting in Artificial Neural Networks* [Master’s thesis, University of Alberta]. <https://doi.org/10.7939/r3-6zv-v5z64>
- Ashley, D. R., Arulkumaran, K., Schmidhuber, J., & Srivastava, R. K. (2022). *Learning relative return policies with upside-down reinforcement learning*. arXiv. <https://doi.org/10.48550/arXiv.2202.12742>

- Ashley, D. R., Chockalingam, V., Kuzma, B., & Bulitko, V. (2019a). Learning to select mates in artificial life. *Proceedings of the Genetic and Evolutionary Computation Conference Companion*, 103–104. <https://doi.org/10.1145/3319619.3322060>
- Ashley, D. R., Chockalingam, V., Kuzma, B., & Bulitko, V. (2019b). Learning to select mates in evolving non-playable characters. *Proceedings of the 2019 IEEE Conference on Games*, 1–8. <https://doi.org/10.1109/CIG.2019.8848114>
- Ashley, D. R., Ghiassian, S., & Sutton, R. S. (2021). *Does the Adam optimizer exacerbate catastrophic forgetting?* arXiv. <https://doi.org/10.48550/arXiv.2102.07686>
- Ashley, D. R., Herrmann, V., Friggstad, Z., Mathewson, K. W., & Schmidhuber, J. (2021). *Automatic embedding of stories into collections of independent media*. arXiv. <https://doi.org/10.48550/arXiv.2111.02216>
- Ashley, D. R., Herrmann, V., Friggstad, Z., & Schmidhuber, J. (2022). On narrative information and the distillation of stories. *Proceedings of the NeurIPS 2022 Workshop on Information-Theoretic Principles in Cognitive Systems*. <https://openreview.net/pdf?id=Q2xpIFubR8M>
- Ashley, D. R., Herrmann, V., Friggstad, Z., & Schmidhuber, J. (2025). On the distillation of stories for transferring narrative arcs in collections of independent media. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47, 697–707. <https://doi.org/10.1109/TPAMI.2024.3480702>
- Ashley, D. R., Kanervisto, A., & Bennett, B. (2021). *Back to square one: Superhuman performance in chutes and ladders through deep neural networks and tree search*. arXiv. <https://doi.org/10.48550/arXiv.2104.00698>
- Ashley, D. R., Lan, G. L., Zhao, C., Dhingra, N., Cai, Z., Chang, E., Zhuge, M., Shi, Y., Chandra, V., & Schmidhuber, J. (2026). *RPRA: predicting an llm-judge for efficient but performant inference*. arXiv. <https://doi.org/10.48550/arXiv.2604.12634>
- Ashley, D. R., Przepióra, J., Chen, Y., Abualsaud, A., Yesmagambet, N., Park, S., Feron, E., & Schmidhuber, J. (2026). *RACAS: controlling diverse robots with a single agentic system*. arXiv. <https://doi.org/10.48550/arXiv.2603.05621>
- Austin, J., Odena, A., Nye, M. I., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C. J., Terry, M., Le, Q. V., & Sutton, C. (2021). *Program synthesis with large language models*. arXiv. <https://arxiv.org/abs/2108.07732>
- Barnes, J. M., & Underwood, B. J. (1959). “Fate” of first-list associations in transfer theory. *Journal of Experimental Psychology*, 58(2), 97–105. <https://doi.org/10.1037/h0047507>
- Bar-On, Y. M., Phillips, R., & Milo, R. (2018). The biomass distribution on earth. *Proceedings of the National Academy of Sciences*, 115(25), 6506–6511. <https://doi.org/10.1073/pnas.1711842115>
- Barrett, B. (2020, February 3). *An artist used 99 phones to fake a google maps traffic jam*. WIRED. <https://www.wired.com/story/99-phones-fake-google-maps-traffic-jam/>
A news article about a person who created a fake traffic jam by carrying a bunch of phones in a wagon around the city.
- Barto, A. G., Sutton, R. S., & Anderson, C. W. (1983). Neuronlike adaptive elements that can solve difficult learning control problems. *IEEE Transactions on Systems, Man, and Cybernetics*, 13(5), 834–846. <https://doi.org/10.1109/TSMC.1983.6313077>
- Bellemare, M. G., Candido, S., Castro, P. S., Gong, J., Machado, M. C., Moitra, S., Ponda, S. S., & Wang, Z. (2020). Autonomous navigation of stratospheric balloons using reinforcement learning. *Nature*, 588(7836), 77–82. <https://doi.org/10.1038/s41586-020-2939-8>
- Bellman, R. E. (1957a). *Dynamic programming*. Princeton University Press.
The original paper that introduced dynamic programming.

- Bellman, R. E. (1957b). A markovian decision process. *Journal of Mathematics and Mechanics*, 6(5), 679–684. <http://www.jstor.org/stable/24900506>
- Betker, J., Goh, G., Jing, L., Brooks, T., Wang, J., Li, L., Ouyang, L., Zhuang, J., Lee, J., Guo, Y., Manassra, W., Dhariwal, P., Chu, C., Jiao, Y., & Ramesh, A. (2023). *Improving image generation with better captions*. OpenAI. <https://cdn.openai.com/papers/dall-e-3.pdf>
- Bittner, R. M., Gu, M., Hernandez, G., Humphrey, E. J., Jehan, T., McCurry, H., & Montecchio, N. (2017). Automatic playlist sequencing and transitions. *Proceedings of the 18th International Society for Music Information Retrieval Conference*, 442–448. <https://archives.ismir.net/ismir2017/paper/000086.pdf>
- BlakeSergin. (2024). *How many r's in strawberry? why is this a very difficult question for the ai?* Reddit. https://www.reddit.com/r/singularity/comments/ienqk04/how_many_rs_in_strawberry_why_is_this_a_very/
- Bogdanov, D., Wack, N., Gómez, E., Gulati, S., Herrera, P., Mayor, O., Roma, G., Salamon, J., Zapata, J. R., & Serra, X. (2013). Essentia: An audio analysis library for music information retrieval. *Proceedings of the 14th International Society for Music Information Retrieval Conference*, 493–498. <https://archives.ismir.net/ismir2013/paper/000177.pdf>
- Bommasani, R., Hudson, D. A., Adeli, E., Altman, R., Arora, S., Arx, S. von, Bernstein, M. S., Bohg, J., Bosselut, A., Brunskill, E., Brynjolfsson, E., Buch, S., Card, D., Castellon, R., Chatterji, N., Chen, A., Creel, K., Davis, J. Q., Demszky, D., ... Liang, P. (2021). *On the opportunities and risks of foundation models*. arXiv. <https://doi.org/10.48550/arXiv.2108.07258>
- Bonds, M. E. (2010). The spatial representation of musical form. *Journal of Musicology*, 27(3), 265–303. <https://doi.org/10.1525/jm.2010.27.3.265>
- Bonnin, G., & Jannach, D. (2013). A comparison of playlist generation strategies for music recommendation and a new baseline scheme. *Papers from the 2013 AAAI Workshop*, 16–23. <https://www.aaai.org/ocs/index.php/WS/AAAIW13/paper/view/7078>
- Bonnin, G., & Jannach, D. (2014). Automated generation of music playlists: Survey and experiments. *ACM Computing Surveys*, 47(2), 26:1–26:35. <https://doi.org/10.1145/2652481>
- Brenneis, D. J. A., Dawson, M. R., & Pilarski, P. M. (2017). Development of the handi hand: An inexpensive, multi-articulating, sensorized hand for machine learning research in myoelectric control. *Proceedings of the 17th Myoelectric Controls Symposium*, 1–4.
- Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., & Zaremba, W. (2016). *OpenAI gym*. arXiv. <https://doi.org/10.48550/arXiv.1606.01540>
- Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D. M., Wu, J., Winter, C., ... Amodei, D. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877–1901. https://proceedings.neurips.cc/paper_files/paper/2020/file/1457cod6bfcb4967418bfb8ac142f64a-Paper.pdf
- Burgard, J. P. (2025). *Sam altman predicts ai will surpass human intelligence by 2030*. Business Insider. <https://www.businessinsider.com/sam-altman-predicts-ai-agi-surpass-human-intelligence-2030-2025-9>
Says Sam Altman predicts AGI will occur before 2030.
- Cacioppo, J. T., & Cacioppo, S. (2014). Social relationships and health: The toxic effects of perceived social isolation. *Social and Personality Psychology Compass*, 8(2), 58–72. <https://doi.org/10.1111/spc3.12087>

- Calisti, M., Giorelli, M., Levy, G., Mazzolai, B., Hochner, B., Laschi, C., & Dario, P. (2011). An octopus-bioinspired solution to movement and manipulation for soft robots. *Bioinspiration & Biomimetics*, 6(3), 036002. <https://doi.org/10.1088/1748-3182/6/3/036002>
- Campbell, J. (2008). *The hero with a thousand faces* (3rd ed.). New World Library.
- Canadian Institute for Health Information. (2019). *National health expenditure trends, 1975 to 2019*. <https://www.cihi.ca/sites/default/files/document/nhex-trends-narrative-report-2019-en-web.pdf>
- Chan, C., Chen, W., Su, Y., Yu, J., Xue, W., Zhang, S., Fu, J., & Liu, Z. (2024). ChatEval: Towards better llm-based evaluators through multi-agent debate. *Proceedings of the 12th International Conference on Learning Representations*. <https://openreview.net/pdf?id=FQepisCUWu>
- Chance, P. (2014). *Learning and behavior* (7th ed.). Cengage Learning.
- Chao, P., Debenedetti, E., Robey, A., Andriushchenko, M., Croce, F., Sehwag, V., Dobriban, E., Flammarion, N., Pappas, G. J., Tramèr, F., Hassani, H., & Wong, E. (2024). *JailbreakBench: An open robustness benchmark for jailbreaking large language models*. arXiv. <https://doi.org/10.48550/arXiv.2404.01318>
Introduced a popular benchmark for testing methods meant to jailbreak LLMs.
- Chao, P., Robey, A., Dobriban, E., Hassani, H., Pappas, G. J., & Wong, E. (2023). *Jailbreaking black box large language models in twenty queries*. arXiv. <https://doi.org/10.48550/arXiv.2310.08419>
Used an LLM and an LLM judge to jailbreak a third LLM.
- Chen, J., Lu, Y., Wang, X., Zeng, H., Huang, J., Gesi, J., Xu, Y., Yao, B., & Wang, D. (2025). *Multi-agent-as-judge: Aligning llm-agent-based automated evaluation with multi-dimensional human evaluation*. arXiv. <https://doi.org/10.48550/arXiv.2507.21028>
- Chen, L., Lu, K., Rajeswaran, A., Lee, K., Grover, A., Laskin, M., Abbeel, P., Srinivas, A., & Mordatch, I. (2021). *Decision transformer: Reinforcement learning via sequence modeling*. arXiv. <https://doi.org/10.48550/arXiv.2106.01345>
- Chen, M., Tworek, J., Jun, H., Yuan, Q., Oliveira Pinto, H. P. de, Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., ... Zaremba, W. (2021). *Evaluating large language models trained on code*. arXiv. <https://arxiv.org/abs/2107.03374>
- Chen, S., Moore, J. L., Turnbull, D. R., & Joachims, T. (2012). Playlist prediction via metric embedding. *Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 714–722. <https://doi.org/10.1145/2339530.2339643>
- Chen, T., Kornblith, S., Norouzi, M., & Hinton, G. E. (2020). A simple framework for contrastive learning of visual representations. *Proceedings of the 37th International Conference on Machine Learning*, 1597–1607. <http://proceedings.mlr.press/v119/chen20j.html>
- Chen, Z., & Liu, B. (2018). *Lifelong machine learning* (2nd ed.). Morgan & Claypool Publishers. <https://doi.org/10.2200/S00832EDrV01Y201802AIM037>
- Chiang, W.-L., Li, Z., Lin, Z., Sheng, Y., Wu, Z., Zhang, H., Zheng, L., Zhuang, S., Zhuang, Y., Gonzalez, J. E., Stoica, I., & Xing, E. P. (2023). *Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality*. LMSYS Corp. <https://lmsys.org/blog/2023-03-30-vicuna/>
Introduced the Vicuna family of small LLM models.
- Cigna. (2018). *Cigna U.S. loneliness index*. <https://www.cigna.com/assets/docs/newsroom/loneliness-survey-2018-full-report.pdf>
- Clarice, L. R.-Y., Shiyong, C., Yang, C. J., & Ren, H. (2023). Biomimetic untethered inflatable origami. In H. Ren (Ed.), *Deployable multimodal machine intelligence: Applications in biomedical*

- engineering* (pp. 123–154). Springer Nature Singapore. https://doi.org/10.1007/978-981-19-5932-5_7
- Clarke, A. C. (1973). *Profiles of the future* (2nd ed.). Harper & Row.
- Cliff, D. (2000). *Hang the DJ: Automatic sequencing and seamless mixing of dance-music tracks*. Hewlett-Packard Company. <https://www.hpl.hp.com/techreports/2000/HPL-2000-104.pdf>
- Cunningham, S. J., Bainbridge, D., & Falconer, A. (2006). 'more of an art than a science': Supporting the creation of playlists and mixes. *Proceedings of the 7th International Society for Music Information Retrieval Conference*, 240–245. <https://archives.ismir.net/ismir2006/paper/000085.pdf>
- Dai, Y., Wang, Y., Ashley, D. R., & Schmidhuber, J. (2026). Efficient morphology-control co-design via stackelberg proximal policy optimization. *Proceedings of the 14th International Conference on Learning Representations*. <https://openreview.net/pdf?id=sJovOOklw>
- Davies, A., & Marshall, A. (2021, September 9). *The WIRED guide to self-driving cars*. WIRED. <https://www.wired.com/story/guide-self-driving-cars/>
- Dayan, P., & Hinton, G. E. (1997). Using expectation-maximization for reinforcement learning. *Neural Computation*, 9(2), 271–278. <https://doi.org/10.1162/neco.1997.9.2.271>
- DeepSeek-AI, Guo, D., Yang, D., Zhang, H., Song, J., Zhang, R., Xu, R., Zhu, Q., Ma, S., Wang, P., Bi, X., Zhang, X., Yu, X., Wu, Y., Wu, Z. F., Gou, Z., Shao, Z., Li, Z., Gao, Z., ... Li, S. S. (2025). *DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning*. arXiv. <https://doi.org/10.48550/arXiv.2501.12948>
- Introduced the DeepSeek family of reasoning LLMs.
- Defferrard, M., Benzi, K., Vandergheynst, P., & Bresson, X. (2017). FMA: a dataset for music analysis. *Proceedings of the 18th International Society for Music Information Retrieval Conference*, 316–323. https://ismir2017.smcnus.org/wp-content/uploads/2017/10/75%5C_Paper.pdf
- Degrave, J., Felici, F., Buchli, J., Neunert, M., Tracey, B., Carpanese, F., Ewalds, T., Hafner, R., Abdolmaleki, A., Casas, D. de las, Donner, C., Fritz, L., Galperti, C., Huber, A., Keeling, J., Tsimpoukelli, M., Kay, J., Merle, A., Moret, J.-M., ... Riedmiller, M. (2022). Magnetic control of tokamak plasmas through deep reinforcement learning. *Nature*, 602(7897), 414–419. <https://doi.org/10.1038/s41586-021-04301-9>
- Degrís, T., Pilarski, P. M., & Sutton, R. S. (2012). Model-free reinforcement learning with continuous action in practice. *Proceedings of the 2012 American Control Conference*, 2177–2182. <https://doi.org/10.1109/ACC.2012.6315022>
- DeJong, G., & Spong, M. W. (1994). Swinging up the acrobot: An example of intelligent control. *Proceedings of the 1994 American Control Conference*, 2, 2158–2162. <https://doi.org/10.1109/ACC.1994.752458>
- Duan, Y., Chen, X., Houthoofd, R., Schulman, J., & Abbeel, P. (2016). Benchmarking deep reinforcement learning for continuous control. *Proceedings of the 33rd International Conference on Machine Learning*, 48, 1329–1338. <http://proceedings.mlr.press/v48/duani16.html>
- Dubois, Y., Galambosi, B., Liang, P., & Hashimoto, T. B. (2024). *Length-controlled AlpacaEval: A simple way to debias automatic evaluators*. arXiv. <https://doi.org/10.48550/arXiv.2404.04475>
- Ebbinghaus, H. (1913). *Memory: A contribution to experimental psychology* (H. A. Ruger & C. E. Bussenius, Trans.). Teachers College Press. (Original work published 1885).
- Edwards, A. L., Dawson, M. R., Hebert, J. S., Sherstan, C., Sutton, R. S., Chan, K. M., & Pilarski, P. M. (2016). Application of real-time machine learning to myoelectric prosthesis control: A case series in adaptive switching. *Prosthetics and Orthotics International*, 40(5), 573–581. <https://doi.org/10.1177/0309364615605373>

- Edwards, A. L., Hebert, J. S., & Pilarski, P. M. (2016). Machine learning and unlearning to autonomously switch between the functions of a myoelectric arm. *Proceedings of the 6th IEEE International Conference on Biomedical Robotics and Biomechatronics*, 514–521. <https://doi.org/10.1109/BIOROB.2016.7523678>
- Edwards, B. (2025). *Anthropic chief says ai could surpass "almost all humans at almost everything" shortly after 2027*. Ars Technica. <https://arstechnica.com/ai/2025/01/anthropic-chief-says-ai-could-surpass-almost-all-humans-at-almost-everything-shortly-after-2027/>
Says that Anthropic predicts AGI will occur in 2027.
- Eerola, T., Lartillot, O., & Toivainen, P. (2009). Prediction of multidimensional emotional ratings in music from audio using multivariate regression models. *Proceedings of the 10th International Society for Music Information Retrieval Conference*, 621–626. <http://ismir2009.ismir.net/proceedings/PS4-8.pdf>
- Eiras, F., Zemour, E., Lin, E., & Mugunthan, V. (2025). *Know thy judge: On the robustness meta-evaluation of LLM safety judges*. arXiv. <https://doi.org/10.48550/arXiv.2503.04474>
- FAANG. (2024). *Advanced prompt engineering: Building production-grade llm systems*. Medium. <https://medium.com/@FAANG/advanced-prompt-engineering-building-production-grade-llm-systems-e2143fd7fffd>
- Faccio, F., Herrmann, V., Ramesh, A., Kirsch, L., & Schmidhuber, J. (2022). *Goal-conditioned generators of deep policies*. arXiv. <https://doi.org/10.48550/arXiv.2207.01570>
- Farquhar, S., & Gal, Y. (2018). *Towards robust evaluations of continual learning*. arXiv. <https://doi.org/10.48550/arXiv.1805.09733>
- Fawzi, A., Balog, M., Huang, A., Hubert, T., Romera-Paredes, B., Barekatin, M., Novikov, A., Ruiz, F. J. R., Schrittwieser, J., Swirszcz, G., Silver, D., Hassabis, D., & Kohli, P. (2022). Discovering faster matrix multiplication algorithms with reinforcement learning. *Nature*, 610(7930), 47–53. <https://doi.org/10.1038/s41586-022-05172-4>
- Fedus, W., Ghosh, D., Martin, J. D., Bellemare, M. G., Bengio, Y., & Larochelle, H. (2020). *On catastrophic interference in atari 2600 games*. arXiv. <https://doi.org/10.48550/arXiv.2002.12499>
- Fernando, C., Banarse, D., Michalewski, H., Osindero, S., & Rocktäschel, T. (2024). Promptbreeder: Self-referential self-improvement via prompt evolution. *Proceedings of the 41st International Conference on Machine Learning*, 235, 13481–13544. <https://raw.githubusercontent.com/mlresearch/v235/main/assets/fernando24a/fernando24a.pdf>
- Frantar, E., Ashkboos, S., Hoefler, T., & Alistarh, D. (2022). *GPTQ: Accurate post-training quantization for generative pre-trained transformers*. arXiv. <https://doi.org/10.48550/arXiv.2210.17323>
- French, R. M. (1991). Using semi-distributed representations to overcome catastrophic forgetting in connectionist networks. *Proceedings of the Thirteenth Annual Conference of the Cognitive Science Society*, 173–178. https://cognitivesciencesociety.org/wp-content/uploads/2019/01/cogsci_13.pdf
- Freytag, G. (1894). *Die technik des dramas*. S. Hirzel.
- Fu, J., Ng, S., Jiang, Z., & Liu, P. (2024). GPTScore: Evaluate as you desire. *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 6556–6576. <https://doi.org/10.18653/v1/2024.naacl-long.365>
- Fu, Y., Wang, X., Tian, Y., & Zhao, J. (2024). *Deep think with confidence*. arXiv. <https://doi.org/10.48550/arXiv.2508.15260>
- Furini, M., & Montangero, M. (2022). Automatic and personalized sequencing of music playlists. *Proceedings of the 2022 IEEE International Conference on Distributed Computing Systems Workshops*, 205–208. <https://doi.org/10.1109/ICDCSW56584.2022.00046>

- Furuta, H., Matsuo, Y., & Gu, S. S. (2021). *Generalized decision transformer for offline hindsight information matching*. arXiv. <https://doi.org/10.48550/arXiv.2111.10364>
- Gabora, L. (2019). Creativity: Linchpin in the quest for a viable theory of cultural evolution. *Current Opinion in Behavioral Sciences*, 27, 77–83. <https://doi.org/10.1016/j.cobeha.2018.09.013>
- Gatys, L. A., Ecker, A. S., & Bethge, M. (2016). Image style transfer using convolutional neural networks. *Proceedings of the 2016 IEEE Conference on Computer Vision and Pattern Recognition*, 2414–2423. <https://doi.org/10.1109/CVPR.2016.265>
- Geramifard, A., Dann, C., Klein, R. H., Dabney, W., & How, J. P. (2015). RLPy: A value-function-based reinforcement learning framework for education and research. *Journal of Machine Learning Research*, 16(46), 1573–1578. <http://jmlr.org/papers/v16/geramifard15a.html>
- Ghiassian, S., Rafiee, B., Lo, Y. L., & White, A. (2020). Improving performance in reinforcement learning by breaking generalization in neural networks. *Proceedings of the 19th International Conference on Autonomous Agents and Multiagent Systems*, 438–446. <http://ifaamas.org/Proceedings/aamas2020/pdfs/p438.pdf>
- Ghiassian, S., Rafiee, B., & Sutton, R. S. (2017). *A first empirical study of emphatic temporal difference learning*. arXiv. <https://doi.org/10.48550/arXiv.1705.04185>
- Ghosh, D., Gupta, A., Reddy, A., Fu, J., Devin, C. M., Eysenbach, B., & Levine, S. (2021). Learning to reach goals via iterated supervised learning. *Proceedings of the 9th International Conference on Learning Representations*. <https://openreview.net/forum?id=rALAoXo6yNJ>
- Glorot, X., & Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. *Proceedings of the 13th International Conference on Artificial Intelligence and Statistics*, 9, 249–256. <http://proceedings.mlr.press/v9/glorot10a/glorot10a.pdf>
- Glorot, X., Bordes, A., & Bengio, Y. (2011). Deep sparse rectifier neural networks. *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics*, 15, 315–323. <http://proceedings.mlr.press/v15/glorot11a/glorot11a.pdf>
- Goodfellow, I. J., Bengio, Y., & Courville, A. C. (2016). *Deep learning*. MIT Press.
- Goodfellow, I. J., Mirza, M., Xiao, D., Courville, A. C., & Bengio, Y. (2013). *An empirical investigation of catastrophic forgetting in gradient-based neural networks*. arXiv. <https://doi.org/10.48550/arXiv.1312.6211>
- Gottschall, J. (2012). *The storytelling animal: How stories make us human*. Houghton Mifflin Harcourt.
- Grattafiori, A., Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Vaughan, A., Yang, A., Fan, A., Goyal, A., Hartshorn, A., Yang, A., Mitra, A., Sravankumar, A., Korenev, A., Hinsvark, A., ... Ma, Z. (2024). *The llama 3 herd of models*. arXiv. <https://doi.org/10.48550/arXiv.2407.21783>
Introduced the Llama 3 family of LLMs.
- Graves, A., & Schmidhuber, J. (2005). Framewise phoneme classification with bidirectional lstm networks. *Proceedings of the 2005 IEEE International Joint Conference on Neural Networks*, 4, 2047–2052. <https://doi.org/10.1109/IJCNN.2005.1556215>
- Grill, J., Strub, F., Altché, F., Tallec, C., Richemond, P. H., Buchatskaya, E., Doersch, C., Pires, B. Á., Guo, Z., Azar, M. G., Piot, B., Kavukcuoglu, K., Munos, R., & Valko, M. (2020). Bootstrap your own latent - A new approach to self-supervised learning. *Advances in Neural Information Processing Systems*, 33, 21271–21284. <https://proceedings.neurips.cc/paper/2020/file/f3ada80d5c4ee70142b17b8192b2958e-Paper.pdf>
- Guo, Q., Wang, R., Guo, J., Li, B., Song, K., Tan, X., Liu, G., Bian, J., & Yang, Y. (2024). Connecting large language models with evolutionary algorithms yields powerful prompt opti-

- mizers. *Proceedings of the 12th International Conference on Learning Representations*. <https://openreview.net/pdf?id=ZG3RaNIso8>
- Gutmann, M., & Hyvärinen, A. (2010). Noise-contrastive estimation: A new estimation principle for unnormalized statistical models. *Proceedings of the 13th International Conference on Artificial Intelligence and Statistics*, 9, 297–304. <http://proceedings.mlr.press/v9/gutmann10a.html>
- Gutstein, S., & Stump, E. (2015). Reduction of catastrophic forgetting with transfer learning and ternary output codes. *Proceedings of the 2015 IEEE International Joint Conference on Neural Networks*, 1–8. <https://doi.org/10.1109/IJCNN.2015.7280416>
- Ha, D., & Schmidhuber, J. (2018). Recurrent world models facilitate policy evolution. *Advances in Neural Information Processing Systems*, 31, 2455–2467. https://proceedings.neurips.cc/paper_files/paper/2018/file/2de5d16682c3c35007e4e92982fia2ba-Paper.pdf
- Hachiya, H., Peters, J., & Sugiyama, M. (2009). Efficient sample reuse in em-based policy search. *Proceedings of the Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, 5781, 469–484. https://doi.org/10.1007/978-3-642-04180-8_48
- Hachiya, H., Peters, J., & Sugiyama, M. (2011). Reward-weighted regression with sample reuse for direct policy search in reinforcement learning. *Neural Computation*, 23(11), 2798–2832. https://doi.org/10.1162/NECO_a_00199
- Hafner, D., Lillicrap, T. P., Ba, J., & Norouzi, M. (2020). Dream to control: Learning behaviors by latent imagination. *Proceedings of the 8th International Conference on Learning Representations*. <https://openreview.net/forum?id=SlOTC4tDS>
Introduced the Dreamer architecture for planning with world models.
- Hafner, D., Lillicrap, T. P., Norouzi, M., & Ba, J. (2021). Mastering atari with discrete world models. *Proceedings of the 9th International Conference on Learning Representations*. <https://openreview.net/forum?id=0oabwyZbOu>
Introduced the DreamerV2 architecture.
- Hafner, D., Pasukonis, J., Ba, J., & Lillicrap, T. P. (2023). *Mastering diverse domains through world models*. arXiv. <https://doi.org/10.48550/arXiv.2301.04104>
Introduced the DreamerV3 architecture.
- Hansen, D. L., & Golbeck, J. (2009). Mixing it up: Recommending collections of items. *Proceedings of the 27th International Conference on Human Factors in Computing Systems*, 1217–1226. <https://doi.org/10.1145/1518701.1518883>
- Hart, P. E., Nilsson, N. J., & Raphael, B. (1968). A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*, 4(2), 100–107. <https://doi.org/10.1109/TSSC.1968.300136>
The original paper for the A* algorithm.
- Hastie, T., Tibshirani, R., & Friedman, J. H. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (2nd ed.). Springer. <https://doi.org/10.1007/978-0-387-84858-7>
- He, K., Zhang, X., Ren, S., & Sun, J. (2015). Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. *Proceedings of the 2015 IEEE International Conference on Computer Vision*, 1026–1034. <https://doi.org/10.1109/ICCV.2015.123>
- He, Z., Cao, P., Wang, C., Jin, Z., Chen, Y., Xu, J., Li, H., Jiang, X., Liu, K., & Zhao, J. (2024). *Agentscourt: Building judicial decision-making agents with court debate simulation and legal knowledge augmentation*. arXiv. <https://doi.org/10.48550/arXiv.2403.02959>
- Herrmann, V., Ashley, D. R., & Schmidhuber, J. (2024). *Automatic album sequencing*. arXiv. <https://doi.org/10.48550/arXiv.2411.07772>

- Hetherington, P. A., & Seidenberg, M. S. (1989). Is there ‘catastrophic interference’ in connectionist networks? *Proceedings of the Eleventh Annual Conference of the Cognitive Science Society*, 26–33. https://cognitivesciencesociety.org/wp-content/uploads/2019/01/cogsci_11.pdf
- Hinton, G. E., Srivastava, N., & Swersky, K. (2012). *RMSProp: Divide the gradient by a running average of its recent magnitude* [PDF slides]. https://www.cs.toronto.edu/~tijmen/csc321/slides/lecture_slides_lec6.pdf
- Hjelm, R. D., Fedorov, A., Lavoie-Marchildon, S., Grewal, K., Trischler, A., & Bengio, Y. (2018). *Learning deep representations by mutual information estimation and maximization*. arXiv. <https://doi.org/10.48550/arXiv.1808.06670>
- Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735–1780. <https://doi.org/10.1162/neco.1997.9.8.1735>
- Hopcroft, J. E., & Karp, R. M. (1973). An $n^{5/2}$ algorithm for maximum matchings in bipartite graphs. *SIAM Journal on Computing*, 2(4), 225–231. <https://doi.org/10.1137/0202019>
- Huang, S., Dossa, R. F. J., Ye, C., Braga, J., Chakraborty, D., Mehta, K., & Araújo, J. G. M. (2022). CleanRL: High-quality single-file implementations of deep reinforcement learning algorithms. *Journal of Machine Learning Research*, 23(274), 1–18. <http://jmlr.org/papers/v23/21-1342.html>
- Huang, Y., Xu, Y. E., Sun, K., Yan, V., Sun, A., Khan, H., Nguyen, J., Kachuee, M., Lin, Z., Liu, Y., Colak, A., Kumar, A., Yih, W., & Dong, X. L. (2025). *Confqa: Answer only if you are confident*. arXiv. <https://doi.org/10.48550/arXiv.2506.07309>
- HuggingFaceH4. (2025). *Aime 2024*. Hugging Face, Inc. https://huggingface.co/datasets/HuggingFaceH4/aime_2024
The AIME 2024 mathematics competition dataset.
- Ikeda, S., Oku, K., & Kawagoe, K. (2009). Analysis of music transition in acoustic feature space for music recommendation. *Proceedings of the 9th International Conference on Machine Learning and Computing*, 77–80. <https://doi.org/10.1145/3055635.3056602>
- Irie, K., Schlag, I., Csordás, R., & Schmidhuber, J. (2022). A modern self-referential weight matrix that learns to modify itself. *Proceedings of the 39th International Conference on Machine Learning*, 162, 9660–9677. <https://proceedings.mlr.press/v162/irie22b/irie22b.pdf>
- Janner, M., Li, Q., & Levine, S. (2021). *Reinforcement learning as one big sequence modeling problem*. arXiv. <https://doi.org/10.48550/arXiv.2106.02039>
- Jarrett, K., Kavukcuoglu, K., Ranzato, M., & LeCun, Y. (2009). What is the best multi-stage architecture for object recognition? *Proceedings of the 2009 IEEE International Conference on Computer Vision*, 2146–2153. <https://doi.org/10.1109/ICCV.2009.5459469>
- Jeong, S., Baek, J., Cho, S., Hwang, S. J., & Park, J. (2024). GPTScore: Evaluate as you desire. *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 7036–7050. <https://doi.org/10.18653/v1/2024-naacl-long.389>
- Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., & Narasimhan, K. R. (2024). SWE-bench: Can language models resolve real-world Github issues? *Proceedings of the 12th International Conference on Learning Representations*. <https://openreview.net/pdf?id=VTF8yNQM66>
- Jin, D., Pan, E., Oufattole, N., Weng, W., Fang, H., & Szolovits, P. (2020). *What disease does this patient have? A large-scale open domain question answering dataset from medical exams*. arXiv. <https://doi.org/10.48550/arXiv.2009.13081>
Introduced the Med QA dataset.
- Jonker, R., & Volgenant, A. (1987). A shortest augmenting path algorithm for dense and sparse linear assignment problems. *Computing*, 38(4), 325–340. <https://doi.org/10.1007/BF02278710>

- Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., Žídek, A., Potapenko, A., Bridgland, A., Meyer, C., Kohl, S. A. A., Ballard, A. J., Cowie, A., Romera-Paredes, B., Nikolov, S., Jain, R., Adler, J., ... Hassabis, D. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873), 583–589. <https://doi.org/10.1038/s41586-021-03819-2>
- Kemker, R., McClure, M., Abitino, A., Hayes, T. L., & Kanan, C. (2018). Measuring catastrophic forgetting in neural networks. *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence*, 3390–3398. <https://aaai.org/ocs/index.php/AAAI/AAAI18/paper/view/16410>
- Kerr, B. (2025). *Creativity*. Encyclopedia Britannica. <https://www.britannica.com/topic/creativity>
- Kingma, D. P., & Ba, J. (2014). *Adam: A method for stochastic optimization*. arXiv. <https://doi.org/10.48550/arXiv.1412.6980>
- Kirkpatrick, J., Pascanu, R., Rabinowitz, N. C., Veness, J., Desjardins, G., Rusu, A. A., Milan, K., Quan, J., Ramalho, T., Grabska-Barwinska, A., Hassabis, D., Clopath, C., Kumaran, D., & Hadsell, R. (2017). Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13), 3521–3526. <https://doi.org/10.1073/pnas.1611835114>
- Kirsch, L., Steenkiste, S. van, & Schmidhuber, J. (2020). Improving generalization in meta reinforcement learning using learned objectives. *Proceedings of the 8th International Conference on Learning Representations*. <https://openreview.net/forum?id=SievHerYPr>
- Kober, J., & Peters, J. (2011). Policy search for motor primitives in robotics. *Machine Learning*, 84(1-2), 171–203. <https://doi.org/10.1007/s10994-010-5223-6>
- Koenig, N. P., & Howard, A. (2004). Design and use paradigms for gazebo, an open-source multi-robot simulator. *Proceedings of the 2004 IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2149–2154. <https://doi.org/10.1109/IROS.2004.1389727>
- Korenkevych, D., Mahmood, A. R., Vasan, G., & Bergstra, J. (2019). Autoregressive policies for continuous control deep reinforcement learning. *Proceedings of the 28th International Joint Conference on Artificial Intelligence*, 2754–2762. <https://doi.org/10.24963/IJCAI.2019/382>
- Kornblith, S., Norouzi, M., Lee, H., & Hinton, G. E. (2019). Similarity of neural network representations revisited. *Proceedings of the 36th International Conference on Machine Learning*, 97, 3519–3529. <http://proceedings.mlr.press/v97/kornblith19a/kornblith19a.pdf>
- Krakauer, D. C. (2023). Unifying complexity science and machine learning. *Frontiers in Complex Systems*, 1, 1–8. <https://doi.org/10.3389/fcpxs.2023.1235202>
- Kumar, A., Peng, X. B., & Levine, S. (2019). *Reward-conditioned policies*. arXiv. <https://doi.org/10.48550/arXiv.1912.13465>
- LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324. <https://doi.org/10.1109/5.726791>
- Lee, K., Nachum, O., Yang, M., Lee, L., Freeman, D., Xu, W., Guadarrama, S., Fischer, I., Jang, E., Michalewski, H., & Mordatch, I. (2022). *Multi-game decision transformers*. arXiv. <https://doi.org/10.48550/arXiv.2205.15241>
- Lee, S., Kim, J., Jun, J., Ha, J., & Zhang, B. (2017). Overcoming catastrophic forgetting by incremental moment matching. *Advances in Neural Information Processing Systems*, 30, 4652–4662. <http://papers.nips.cc/paper/7051-overcoming-catastrophic-forgetting-by-incremental-moment-matching.pdf>
- Levenshtein, V. (1966). Binary codes capable of correcting deletions, insertions, and reversals. *Doklady Akademii Nauk SSSR*, 163(4), 845–848
- Introduced the Levenshtein distance metric.

- Leviathan, Y., Kalman, M., & Matias, Y. (2023). Fast inference from transformers via speculative decoding. *Proceedings of the 40th International Conference on Machine Learning*, 202, 19274–19286. <https://proceedings.mlr.press/v202/leviathan23a/leviathan23a.pdf>
- Levine, S., Pastor, P., Krizhevsky, A., Ibarz, J., & Quillen, D. (2018). Learning hand-eye coordination for robotic grasping with deep learning and large-scale data collection. *International Journal of Robotics Research*, 37(4-5), 421–436. <https://doi.org/10.1177/0278364917710318>
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474. https://proceedings.neurips.cc/paper_files/paper/2020/file/6b493230205f780e1bc26945df7481e5-Paper.pdf
- Introduced the RAG system, which is used extensively with LLMs.
- Li, B., Wu, W., Tang, Z., Shi, L., Yang, J., Li, J., Yao, S., Qian, C., Hui, B., Zhang, Q., Yu, Z., Du, H., Yang, P., Lin, D., Peng, C., & Chen, K. (2024). *DevBench: A comprehensive benchmark for software development*. arXiv. <https://doi.org/10.48550/arXiv.2403.08604>
- Li, H., Dong, Q., Chen, J., Su, H., Zhou, Y., Ai, Q., Ye, Z., & Liu, Y. (2024). *LLMs-as-Judges: A comprehensive survey on llm-based evaluation methods*. arXiv. <https://doi.org/10.48550/arXiv.2412.05579>
- Li, H., Li, Y., Tian, A., Tang, T., Xu, Z., Chen, X., Hu, N., Dong, W., Li, Q., & Chen, L. (2025). A survey on large language model acceleration based on KV cache management. *Transactions on Machine Learning Research*. <https://openreview.net/pdf?id=z3JZzu9EA3>
- Li, L., Chai, Y., Wang, S., Sun, Y., Tian, H., Zhang, N., & Wu, H. (2024). Tool-augmented reward modeling. *Proceedings of the 12th International Conference on Learning Representations*. <https://openreview.net/pdf?id=d94xogWTUX>
- Li, S., Xu, C., Wang, J., Gong, X., Chen, C., Zhang, J., Wang, J., Lam, K., & Ji, S. (2025). *Llms cannot reliably judge (yet?): A comprehensive assessment on the robustness of llm-as-a-judge*. arXiv. <https://doi.org/10.48550/arXiv.2506.09443>
- Liang, T., He, Z., Jiao, W., Wang, X., Wang, Y., Wang, R., Yang, Y., Shi, S., & Tu, Z. (2024). Encouraging divergent thinking in large language models through multi-agent debate. *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 17889–17904. <https://doi.org/10.18653/v1/2024.emnlp-main.992>
- Lin, L. (1992). Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine Learning*, 8(3-4), 293–321. <https://doi.org/10.1007/BF00992699>
- Littman, M. L., Sutton, R. S., & Singh, S. (2001). Predictive representations of state. *Advances in Neural Information Processing Systems*, 14, 1555–1561. <https://proceedings.neurips.cc/paper/2001/file/1e4d36177d71bbb3558e43af9577d70e-Paper.pdf>
- Liu, H., Ning, R., Teng, Z., Liu, J., Zhou, Q., & Zhang, Y. (2023). *Evaluating the logical reasoning ability of ChatGPT and GPT-4*. arXiv. <https://doi.org/10.48550/arXiv.2304.03439>
- Liu, H., Wang, Z., Chen, X., Li, Z., Xiong, F., Yu, Q., & Zhang, W. (2025). HopRAG: Multi-hop reasoning for logic-aware retrieval-augmented generation. *Findings of the Association for Computational Linguistics*, 1897–1913. <https://doi.org/10.18653/v1/2025.findings-acl.97>
- Liu, V. (2019). *Sparse Representation Neural Networks for Online Reinforcement Learning* [Master’s thesis, University of Alberta]. <https://doi.org/10.7939/r3-9f56-q561>
- Liu, V., Kumaraswamy, R., Le, L., & White, M. (2019). The utility of sparse representations for control in reinforcement learning. *Proceedings of the Thirty-Third AAAI Conference on Artificial Intelligence*, 4384–4391. <https://doi.org/10.1609/aaai.v33i01.33014384>

- Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., Zhang, S., Deng, X., Zeng, A., Du, Z., Zhang, C., Shen, S., Zhang, T., Su, Y., Sun, H., ... Tang, J. (2024). AgentBench: Evaluating LLMs as agents. *Proceedings of the 12th International Conference on Learning Representations*. <https://openreview.net/pdf?id=zAdUBoaCTQ>
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., & Zhu, C. (2023). G-Eval: NLG evaluation using Gpt-4 with better human alignment. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2511–2522. <https://doi.org/10.18653/v1/2023.emnlp-main.153>
- Liu, Z., Zhao, C., Iandola, F. N., Lai, C., Tian, Y., Fedorov, I., Xiong, Y., Chang, E., Shi, Y., Krishnamoorthi, R., Lai, L., & Chandra, V. (2024). MobileLLM: Optimizing sub-billion parameter language models for on-device use cases. *Proceedings of the 41st International Conference on Machine Learning*, 235, 32431–32454. <https://raw.githubusercontent.com/mlresearch/v235/main/assets/liu24ce/liu24ce.pdf>
- Löwe, S., O'Connor, P., & Veeling, B. S. (2019). Putting an end to end-to-end: Gradient-isolated learning of representations. *Advances in Neural Information Processing Systems*, 32, 3033–3045. <https://proceedings.neurips.cc/paper/2019/file/851300ee84c2b80ed40f51ed26d866fc-Paper.pdf>
- Lucchi, M., Zindler, F., Mühlbacher-Karrer, S., & Pichler, H. (2020). robo-gym - an open source toolkit for distributed deep reinforcement learning on real and simulated robots. *Proceedings of the 2020 IEEE/RSJ International Conference on Intelligent Robots and Systems*, 5364–5371. <https://doi.org/10.1109/IROS45743.2020.9340956>
- Ma, C., Ashley, D. R., Wen, J., & Bengio, Y. (2020). *Universal successor features for transfer reinforcement learning*. arXiv. <https://doi.org/10.48550/arXiv.2001.04025>
- Ma, R. R., Odhner, L., & Dollar, A. M. (2013). A modular, open-source 3D printed underactuated hand. *Proceedings of the IEEE International Conference on Robotics and Automation*, 2737–2743. <https://doi.org/10.1109/ICRA.2013.6630954>
- Maillet, F., Eck, D., Desjardins, G., & Lamere, P. (2009). Steerable playlist generation by learning song similarity from radio station playlists. *Proceedings of the 10th International Society for Music Information Retrieval Conference*, 345–350. <https://archives.ismir.net/ismir2009/paper/000012.pdf>
- Mao, X., Xu, H., Li, F., Jin, Z., Chen, W., Zhang, W., & Luu, A. T. (2025). As simple as fine-tuning: LLM alignment via bidirectional negative feedback loss. *Proceedings of the 13th International Conference on Learning Representations*. <https://openreview.net/pdf?id=fsX9nFwMNj>
- Mar, R. A. (2011). The neural bases of social cognition and story comprehension. *Annual review of psychology*, 62, 103–134. <https://doi.org/10.1146/annurev-psych-120709-145406>
- Marcelo, P. (2023, May 23). *Fact focus: Fake image of pentagon explosion briefly sends jitters through stock market*. Associated Press. <https://apnews.com/article/pentagon-explosion-misinformation-stock-market-ai-96f534c790872fde67012ee81b5ed6a4>
- Martinez, R. V., Glavan, A. C., Keplinger, C., Oyetibo, A. I., & Whitesides, G. M. (2014). Soft actuators and robots that are resistant to mechanical damage. *Advanced Functional Materials*, 24(20), 3003–3010. <https://doi.org/10.1002/adfm.201303676>
- Masse, N. Y., Grant, G. D., & Freedman, D. J. (2018). Alleviating catastrophic forgetting using context-dependent gating and synaptic stabilization. *Proceedings of the National Academy of Sciences*, 115(44), E10467–E10475. <https://doi.org/10.1073/pnas.1803839115>
- Mathewson, K. W., Castro, P. S., Cherry, C., Foster, G. F., & Bellemare, M. G. (2020). Shaping the narrative arc: Information-theoretic collaborative dialogue. *Proceedings of the 11th International*

- Conference on Computational Creativity*, 9–16. <http://computationalcreativity.net/iccc20/papers/010-iccc20.pdf>
- Matusik, H., Liu, C., & Rus, D. (2024). Directly 3d printed, pneumatically actuated multi-material robotic hand. *Proceedings of the IEEE International Conference on Robotics and Automation*, 2986–2992. <https://doi.org/10.1109/ICRA57147.2024.10610016>
- McCarthy, J. (2007). *What is artificial intelligence?* <http://www-formal.stanford.edu/jmc/whatisai.pdf>
- McCloskey, M., & Cohen, N. J. (1989). Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 24, 109–165. [https://doi.org/10.1016/S0079-7421\(08\)60536-8](https://doi.org/10.1016/S0079-7421(08)60536-8)
- McGeoch, J. A. (1932). Forgetting and the law of disuse. *Psychological Review*, 39(4), 352–370. <https://doi.org/10.1037/h0069819>
- Meta AI. (2024). *Llama 3.2: Revolutionizing edge ai and vision with open, customizable models*. Meta Platforms, Inc. <https://llama.com/blog/llama-3-2-vision-llms-edge-ai>
- Meta AI. (2025). *The llama 4 herd: The beginning of a new era of natively multimodal ai innovation*. Meta Platforms, Inc. <https://llama.com/blog/llama-4-multimodal-intelligence>
- Minsky, M. (1986). *The society of mind*. Simon & Schuster.
- Mirhoseini, A., Goldie, A., Yazgan, M., Jiang, J. W., Songhori, E., Wang, S., Lee, Y., Johnson, E., Pathak, O., Nazi, A., Pak, J., Tong, A., Srinivasa, K., Hang, W., Tuncer, E., Le, Q. V., Laudon, J., Ho, R., Carpenter, R., & Dean, J. (2021). A graph placement methodology for fast chip design. *Nature*, 594(7862), 207–212. <https://doi.org/10.1038/s41586-021-03544-w>
- Mirzadeh, S., Farajtabar, M., Pascanu, R., & Ghasemzadeh, H. (2020). Understanding the role of training regimes in continual learning. *Advances in Neural Information Processing Systems*, 33, 7308–7320. <https://proceedings.neurips.cc/paper/2020/file/518a38cc9a0173dob2dco88166981cf8-Paper.pdf>
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M. A., Fidjeland, A., Ostrovski, G., Petersen, S., Beattie, C., Sadik, A., Antonoglou, I., King, H., Kumaran, D., Wierstra, D., Legg, S., & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529–533. <https://doi.org/10.1038/nature14236>
- Modayil, J., & Sutton, R. S. (2014). Prediction driven behavior: Learning predictions that drive fixed responses. *The AAAI-14 Workshop on Artificial Intelligence and Robotics*. <http://incompleteideas.net/papers/MS-AAAIws-2014.pdf>
- Modayil, J., White, A., & Sutton, R. S. (2014). Multi-timescale nexting in a reinforcement learning robot. *Adaptive Behavior*, 22(2), 146–160. <https://doi.org/10.1177/1059712313511648>
- Montgomery, D. O. (2002). *The rock concept album: Context and analysis* [Doctoral dissertation, University of Toronto]. <https://hdl.handle.net/1807/119202>
- Moore, A. W. (1990). *Efficient memory-based learning for robot control* [Doctoral dissertation, University of Cambridge]. <https://www.cl.cam.ac.uk/techreports/UCAM-CL-TR-209.pdf>
- Moruzzi, C. (2021). Measuring creativity: An account of natural and artificial creativity. *European Journal for Philosophy of Science*, 11(1). <https://doi.org/10.1007/s13194-020-00313-w>
- Muldrew, W., Hayes, P., Zhang, M., & Barber, D. (2024). Active preference learning for large language models. *Proceedings of the 41st International Conference on Machine Learning*, 36577–36590. <https://openreview.net/pdf?id=CTgEV6qgUy>
- Nair, V., & Hinton, G. E. (2010). Rectified linear units improve restricted boltzmann machines. *Proceedings of the 27th International Conference on Machine Learning*, 807–814. <http://www.icml2010.org/papers/432.pdf>

- National Highway Traffic Safety Administration. (2019). *2017 quick facts*. <https://crashstats.nhtsa.dot.gov/Api/Public/ViewPublication/812747>
- Neto, P. A. S. O., Hartmann, M., Luck, G., & Toiviainen, P. (2024). The algorithmic nature of song-sequencing: Statistical regularities in music albums. *Journal of New Music Research*, 1–15. <https://doi.org/10.1080/09298215.2024.2423610>
- Neumann, J. von. (1958). *The computer and the brain* (1st ed.). Yale University Press.
- Ng, A. Y., Coates, A., Diel, M., Ganapathi, V., Schulte, J., Tse, B., Berger, E., & Liang, E. (2006). Autonomous inverted helicopter flight via reinforcement learning. *Proceedings of the 9th International Symposium on Experimental Robotics*, 363–372. https://doi.org/10.1007/11552246_35
- Noffsinger, J., Patel, M., Sachdeva, P., Bhan, A., Chang, H., & Goodpaster, M. (2025). The cost of compute: A \$7 trillion dollar race to scale data centers. *McKinsey Quarterly*. <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/the-cost-of-compute-a-7-trillion-dollar-race-to-scale-data-centers>
Quantifies the large amount of investment being spent in the AI race.
- Olson, E. (2011). AprilTag: A robust and flexible visual fiducial system. *Proceedings of the IEEE International Conference on Robotics and Automation*, 3400–3407. <https://doi.org/10.1109/ICRA.2011.5979561>
- Oord, A. van den, Dieleman, S., & Schrauwen, B. (2013). Deep content-based music recommendation. *Advances in Neural Information Processing Systems*, 26, 2643–2651. https://proceedings.neurips.cc/paper_files/paper/2013/file/b3ba8f1bee1238a2f37603d90b58898d-Paper.pdf
- Oord, A. van den, Li, Y., & Vinyals, O. (2018). *Representation learning with contrastive predictive coding*. arXiv. <https://doi.org/10.48550/arXiv.1807.03748>
Introduced contrastive learning.
- OpenAI. (2022). *Introducing ChatGPT*. <https://openai.com/blog/chatgpt>
- OpenAI. (2025). *Introducing gpt-oss*. <https://openai.com/index/introducing-gpt-oss/>
- Oramas, S., Anke, L. E., Lawlor, A., Serra, X., & Saggion, H. (2016). Exploring customer reviews for music genre classification and evolutionary studies. *Proceedings of the 17th International Society for Music Information Retrieval Conference*, 150–156. <https://archives.ismir.net/ismir2016/paper/000240.pdf>
- Oudeyer, P., Kaplan, F., & Hafner, V. V. (2007). Intrinsic motivation systems for autonomous mental development. *IEEE Transactions on Evolutionary Computation*, 11(2), 265–286. <https://doi.org/10.1109/TEVC.2006.890271>
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askell, A., Welinder, P., Christiano, P. F., Leike, J., & Lowe, R. (2022). Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35, 27730–27744. https://papers.nips.cc/paper_files/paper/2022/file/b1efde53be364a73914f58805a001731-Paper-Conference.pdf
Introduced InstructGPT, the first instruct model and the precursor to ChatGPT.
- Padigela, H., Shah, C., & Juyal, D. (2025). *ML-Dev-Bench: Comparative analysis of AI agents on ML development workflows*. arXiv. <https://doi.org/10.48550/arXiv.2502.00964>
- Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., & Bernstein, M. S. (2023). *Generative agents: Interactive simulacra of human behavior*. arXiv. <https://doi.org/10.48550/arXiv.2304.03442>

- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang, E. Z., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). Pytorch: An imperative style, high-performance deep learning library. *Advances in Neural Information Processing Systems*, 32, 8024–8035. https://proceedings.neurips.cc/paper_files/paper/2019/file/bdbca288fee7f92f2bfa9f7012727740-Paper.pdf
- Pecher, B., Srba, I., & Bielikova, M. (2022). *Comparing specialised small and general large language models on text classification: 100 labelled samples to achieve break-even performance*. arXiv. <https://doi.org/10.48550/arXiv.2402.12819>
- Peng, X. B., Kumar, A., Zhang, G., & Levine, S. (2019). *Advantage-weighted regression: Simple and scalable off-policy reinforcement learning*. arXiv. <https://doi.org/10.48550/arXiv.1910.00177>
- Peters, J., & Schaal, S. (2007a). Applying the episodic natural actor-critic architecture to motor primitive learning. *Proceedings of the 15th European Symposium on Artificial Neural Networks*, 295–300. <https://www.esann.org/sites/default/files/proceedings/legacy/es2007-125.pdf>
- Peters, J., & Schaal, S. (2007b). Reinforcement learning by reward-weighted regression for operational space control. *Proceedings of the 24th International Conference on Machine Learning*, 227, 745–750. <https://doi.org/10.1145/1273496.1273590>
- Peters, J., & Schaal, S. (2008). Natural actor-critic. *Neurocomputing*, 71(7-9), 1180–1190. <https://doi.org/10.1016/j.neucom.2007.11.026>
- Pilarski, P. M., Dawson, M. R., Degris, T., Carey, J. P., Chan, K. M., Hebert, J. S., & Sutton, R. S. (2013). Adaptive artificial limbs: A real-time approach to prediction and anticipation. *IEEE Robotics & Automation Magazine*, 20(1), 53–64. <https://doi.org/10.1109/MRA.2012.2229948>
- Pilarski, P. M., Dawson, M. R., Degris, T., Carey, J. P., & Sutton, R. S. (2012). Dynamic switching and real-time machine learning for improved human control of assistive biomedical robots. *Proceedings of the 4th IEEE International Conference on Biomedical Robotics and Biomechatronics*, 296–302. <https://doi.org/10.1109/BioRob.2012.6290309>
- Pilarski, P. M., Dick, T. B., & Sutton, R. S. (2013). Real-time prediction learning for the simultaneous actuation of multiple prosthetic joints. *Proceedings of the 13th IEEE International Conference on Rehabilitation Robotics*, 1–8. <https://doi.org/10.1109/ICORR.2013.6650435>
- Project Apertus. (2025). *Apertus: Democratizing open and compliant llms for global language environments*. Swiss AI Initiative. https://github.com/swiss-ai/apertus-tech-report/blob/main/Apertus_Tech_Report.pdf
- Puterman, M. L. (1994). *Markov decision processes: Discrete stochastic dynamic programming*. Wiley. <https://doi.org/10.1002/9780470316887>
- Qian, N. (1999). On the momentum term in gradient descent learning algorithms. *Neural Networks*, 12(1), 145–151. [https://doi.org/10.1016/S0893-6080\(98\)00116-6](https://doi.org/10.1016/S0893-6080(98)00116-6)
- Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., Lin, Y., Cong, X., Tang, X., Qian, B., Zhao, S., Hong, L., Tian, R., Xie, R., Zhou, J., Gerstein, M., Li, D., Liu, Z., & Sun, M. (2024). ToolLLM: Facilitating large language models to master 16000+ real-world APIs. *Proceedings of the 12th International Conference on Learning Representations*. <https://openreview.net/pdf?id=dHng2OoJjr>
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I. (2021). Learning transferable visual models from natural language supervision. *Proceedings of the 38th International Conference on Machine Learning*, 139, 8748–8763. <http://proceedings.mlr.press/v139/radford21a/radford21a.pdf>

- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). *Language models are unsupervised multitask learners*. OpenAI. https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf
- Ramesh, A., Dhariwal, P., Nichol, A., Chu, C., & Chen, M. (2022). *Hierarchical text-conditional image generation with CLIP latents*. arXiv. <https://doi.org/10.48550/arXiv.2204.06125>
- Ramesh, A., Pavlov, M., Goh, G., Gray, S., Voss, C., Radford, A., Chen, M., & Sutskever, I. (2021). Zero-shot text-to-image generation. *Proceedings of the 38th International Conference on Machine Learning*, 139, 8821–8831. <http://proceedings.mlr.press/v139/ramesh21a.html>
- Ratcliff, R. (1990). Connectionist models of recognition memory: Constraints imposed by learning and forgetting functions. *Psychological Review*, 97(2), 285–308. <https://doi.org/10.1037/0033-295X.97.2.285>
- Reagan, A. J., Mitchell, L., Kiley, D., Danforth, C. M., & Dodds, P. S. (2016). The emotional arcs of stories are dominated by six basic shapes. *EPJ Data Science*, 5(1), 31–42. <https://doi.org/10.1140/epjds/s13688-016-0093-1>
- Reicha, A. (1826). *Traité de haute composition musicale* (Vol. 2). Richault.
- Reid, E. (1981). Negative transfer and proactive interference: Some confusion in introductory psychology texts. *Teaching of Psychology*, 8(2), 109–110. https://doi.org/10.1207/s15328023top0802_15
- Ren, J., Zhao, Y., Vu, T., Liu, P. J., & Lakshminarayanan, B. (2023). Self-evaluation improves selective generation in large language models. *Proceedings on "I Can't Believe It's Not Better: Failure Modes in the Age of Foundation Models" at NeurIPS 2023 Workshops*, 49–64. <https://proceedings.mlr.press/v239/ren23a/ren23a.pdf>
- RIAA. (2021). *Gold & platinum*. Recording Industry Association of America. https://www.riaa.com/gold-platinum/?tab_active=awards_by_album
- Ricoeur, P. (1991). Narrative identity. *Philosophy today*, 35(1), 73–81. <https://doi.org/10.5840/philtoday199135136>
- Riemer, M., Cases, I., Ajemian, R., Liu, M., Rish, I., Tu, Y., & Tesauero, G. (2019). Learning to learn without forgetting by maximizing transfer and minimizing interference. *Proceedings of the 7th International Conference on Learning Representations*. <https://openreview.net/forum?id=BigTShAct7>
- Ring, M. B. (1997). CHILD: A first step towards continual learning. *Machine Learning*, 28(1), 77–104. <https://doi.org/10.1023/A:1007331723572>
- Ring, M. B. (2021). *Representing knowledge as predictions (and state as knowledge)*. arXiv. <https://doi.org/10.48550/arXiv.2112.06336>
- Roth, E. (2025). *Openai says chatgpt users send over 2.5 billion prompts every day*. The Verge. <https://www.theverge.com/news/710867/openai-chatgpt-daily-prompts-2-billion>
A news article about the number of people using AI nowadays, with OpenAI processing more than 2.5 billion prompts, but only about 330 million coming from US users.
- Roy, R., Raiman, J., Kant, N., Elkin, I., Kirby, R., Siu, M. Y., Oberman, S. F., Godil, S., & Catanzaro, B. (2021). PrefixRL: Optimization of parallel prefix circuits using deep reinforcement learning. *Proceedings of the 58th ACM/IEEE Design Automation Conference*, 853–858. <https://doi.org/10.1109/DAC18074.2021.9586094>
- Rubio, F., Valero, F., & Llopis-Albert, C. (2019). A review of mobile robots: Concepts, methods, theoretical framework, and applications. *International Journal of Advanced Robotic Systems*, 16(2), 1–22. <https://doi.org/10.1177/1729881419839596>
- Ruder, S. (2016). *An overview of gradient descent optimization algorithms*. arXiv. <https://doi.org/10.48550/arXiv.1609.04747>

- Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323(6088), 533–536. <https://doi.org/10.1038/323533a0>
- Rus, D., & Tolley, M. T. (2015). Design, fabrication and control of soft robots. *Nature*, 521(7553), 467–475. <https://doi.org/10.1038/nature14543>
- Russell, S. J., & Norvig, P. (2003). *Artificial intelligence - a modern approach* (2nd ed.). Prentice Hall.
- Saeed, A., Grangier, D., & Zeghidour, N. (2021). Contrastive learning of general-purpose audio representations. *Proceedings of the 2021 IEEE International Conference on Acoustics, Speech and Signal Processing*, 3875–3879. <https://doi.org/10.1109/ICASSP39728.2021.9413528>
- Saleh, J. H., Mark, G., & Jordan, N. C. (2009). Flexibility: A multi-disciplinary literature review and a research agenda for designing flexible engineering systems. *Journal of engineering design*, 20(3), 307–323. <https://doi.org/10.1080/09544820701870813>
- Salustowicz, R., Wiering, M. A., & Schmidhuber, J. (1997). On learning soccer strategies. *Proceedings of the 7th International Conference on Artificial Neural Networks*, 1327, 769–774. <https://doi.org/10.1007/BFBo020247>
- Sarroff, A. M., & Casey, M. (2012). Modeling and predicting song adjacencies in commercial albums. *Proceedings of the 9th Sound and Music Computing Conference*, 364–371. <http://smc.afim-asso.org/smc12/smcnetwork.org/node/1676.html>
- Satopaa, V., Albrecht, J. R., Irwin, D. E., & Raghavan, B. (2011). Finding a "kneedle" in a haystack: Detecting knee points in system behavior. *Proceedings of the 2011 IEEE International Conference on Distributed Computing Systems Workshops*, 166–171. <https://doi.org/10.1109/ICDCSW.2011.20>
Proposed an algorithm to automatically find the knee of a curve.
- Schaul, T., Horgan, D., Gregor, K., & Silver, D. (2015). Universal value function approximators. *Proceedings of the 32nd International Conference on Machine Learning*, 37, 1312–1320. <https://proceedings.mlr.press/v37/schaul15.pdf>
- Schaul, T., Pape, L., Glasmachers, T., Graziano, V., & Schmidhuber, J. (2011). Coherence progress: A measure of interestingness based on fixed compressors. *Proceedings of the 4th International Conference on Artificial General Intelligence*, 21–30. https://doi.org/10.1007/978-3-642-22887-2_3
- Schaul, T., & Schmidhuber, J. (2008). A scalable neural network architecture for board games. *Proceedings of the 2008 IEEE Symposium on Computational Intelligence and Games*, 357–364. <https://doi.org/10.1109/CIG.2008.5035662>
- Schick, T., Dwivedi-Yu, J., Dessi, R., Raileanu, R., Lomeli, M., Hambro, E., Zettlemoyer, L., Cancedda, N., & Scialom, T. (2023). Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems*, 36, 68539–68551. https://proceedings.neurips.cc/paper_files/paper/2023/file/d842425e4bf79ba039352da0f658a906-Paper-Conference.pdf
- Schlag, I., Irie, K., & Schmidhuber, J. (2021). Linear transformers are secretly fast weight programmers. *Proceedings of the 38th International Conference on Machine Learning*, 9355–9366. <http://proceedings.mlr.press/v139/schlag21a.html>
Showed the formal equivalence between linear transformers and fast weight programmers.
- Schmidhuber, J. (1991). Curious model-building control systems. *Proceedings of the 1991 IEEE International Joint Conference on Neural Networks*, 2, 1458–1463. <https://doi.org/10.1109/IJCNN.1991.170605>
- Schmidhuber, J. (1992). Learning to control fast-weight memories: An alternative to recurrent nets. *Neural Computation*, 4(1), 131–139. <https://doi.org/10.1162/neco.1992.4.1.131>

- Schmidhuber, J. (2006). Developmental robotics, optimal artificial curiosity, creativity, music, and the fine arts. *Connection Science*, 18(2), 173–187. <https://doi.org/10.1080/09540090600768658>
- Schmidhuber, J. (2007). Simple algorithmic principles of discovery, subjective beauty, selective attention, curiosity & creativity. *Proceedings of the 10th International Conference on Discovery Science*, 4755, 26–38. https://doi.org/10.1007/978-3-540-75488-6_3
- Schmidhuber, J. (2008). *Driven by compression progress: A simple principle explains essential aspects of subjective beauty, novelty, surprise, interestingness, attention, curiosity, creativity, art, science, music, jokes*. arXiv. <https://doi.org/10.48550/arXiv.0812.4360>
- Schmidhuber, J. (2015). *On learning to think: Algorithmic information theory for novel combinations of reinforcement learning controllers and recurrent neural world models*. arXiv. <https://doi.org/10.48550/arXiv.1511.09249>
- Schmidhuber, J. (2018). *One big net for everything*. arXiv. <https://doi.org/10.48550/arXiv.1802.08864>
- Schmidhuber, J. (2019). *Reinforcement learning upside down: Don't predict rewards - just map them to actions*. arXiv. <https://doi.org/10.48550/arXiv.1912.02875>
- Schreiber, H., & Müller, M. (2019). *Musical tempo and key estimation using convolutional neural networks with directional filters*. arXiv. <https://doi.org/10.48550/arXiv.1903.10839>
- Schubert, E. (1999). Measuring emotion continuously: Validity and reliability of the two-dimensional emotion-space. *Australian Journal of Psychology*, 51(3), 154–165. <https://doi.org/10.1080/00049539908255353>
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). *Proximal policy optimization algorithms*. arXiv. <https://doi.org/10.48550/arXiv.1707.06347>
- Schuermans, D. (2023). *Memory augmented large language models are computationally universal*. arXiv. <https://doi.org/10.48550/arXiv.2301.04589>
- Senior, A. W., Evans, R., Jumper, J., Kirkpatrick, J., Sifre, L., Green, T., Qin, C., Židek, A., Nelson, A. W. R., Bridgland, A., Penedones, H., Petersen, S., Simonyan, K., Crossan, S., Kohli, P., Jones, D. T., Silver, D., Kavukcuoglu, K., & Hassabis, D. (2020). Improved protein structure prediction using potentials from deep learning. *Nature*, 577(7792), 706–710. <https://doi.org/10.1038/s41586-019-1923-7>
- Sherstan, C., Ashley, D. R., Bennett, B., Young, K., White, A., White, M., & Sutton, R. S. (2018). Comparing direct and indirect temporal-difference methods for estimating the variance of the return. *Proceedings of the 34th Conference on Uncertainty in Artificial Intelligence*, 63–72. <http://auai.org/uai2018/proceedings/papers/35.pdf>
- Sherstan, C., Bennett, B., Young, K., Ashley, D. R., White, A., White, M., & Sutton, R. S. (2018). *Directly estimating the variance of the λ -return using temporal-difference methods*. arXiv. <https://doi.org/10.48550/arXiv.1801.08287>
- Shi, L., Zhang, H., Yao, Y., Li, Z., & Zhao, H. (2024). *Keep the cost down: A review on methods to optimize llm's KV-cache consumption*. arXiv. <https://doi.org/10.48550/arXiv.2407.18003>
- Shin, T., Razeghi, Y., IV, R. L. L., Wallace, E., & Singh, S. (2020). AutoPrompt: Eliciting knowledge from language models with automatically generated prompts. *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, 4222–4235. <https://doi.org/10.18653/v1/2020.emnlp-main.346>
- Siegenfeld, A. F., & Bar-Yam, Y. (2020). An introduction to complex systems science and its applications. *Complexity*, 2020, 1–16. <https://doi.org/10.1155/2020/6105872>
- Sieker, J., Junker, S., Utescher, R., Attari, N., Wersing, H., Buschmeier, H., & Zarriß, S. (2024). The illusion of competence: Evaluating the effect of explanations on users' mental models of visual question answering systems. *Proceedings of the 2024 Conference on Empirical Methods*

- in *Natural Language Processing*, 19459–19475. <https://doi.org/10.18653/V1/2024.EMNLP-MAIN.1084>
- Silva, B. C. da, Konidaris, G. D., & Barto, A. G. (2012). Learning parameterized skills. *Proceedings of the 29th International Conference on Machine Learning*, 1679–1686. <http://icml.cc/2012/papers/826.pdf>
- Silver, D. L., Yang, Q., & Li, L. (2013). Lifelong machine learning systems: Beyond learning algorithms. *Lifelong Machine Learning: Papers from the 2013 AAAI Spring Symposium*, 49–55. <https://www.aaai.org/ocs/index.php/SSS/SSS13/paper/view/5802>
- Silver, D., Hasselt, H. van, Hessel, M., Schaul, T., Guez, A., Harley, T., Dulac-Arnold, G., Reichert, D. P., Rabinowitz, N. C., Barreto, A., & Degris, T. (2017). The predictron: End-to-end learning and planning. *Proceedings of the 34th International Conference on Machine Learning*, 70, 3191–3199. <http://proceedings.mlr.press/v70/silver17a/silver17a.pdf>
Introduced the Predictron architecture for planning in reinforcement learning.
- Silver, D., Huang, A., Maddison, C. J., Guez, A., Sifre, L., Driessche, G. van den, Schrittwieser, J., Antonoglou, I., Panneershelvam, V., Lanctot, M., Dieleman, S., Grewe, D., Nham, J., Kalchbrenner, N., Sutskever, I., Lillicrap, T. P., Leach, M., Kavukcuoglu, K., Graepel, T., & Hassabis, D. (2016). Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587), 484–489. <https://doi.org/10.1038/nature16961>
- Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., Lillicrap, T., Simonyan, K., & Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, 362(6419), 1140–1144. <https://doi.org/10.1126/science.aar6404>
- Skinner, B. F. (1950). Are theories of learning necessary? *Psychological Review*, 57(4), 193–216. <https://doi.org/10.1037/h0054367>
- Sodhani, S., Chandar, S., & Bengio, Y. (2020). Toward training recurrent neural networks for lifelong learning. *Neural Computation*, 32(1), 1–35. https://doi.org/10.1162/neco_a_01246
- Soldan, M., Pardo, A., Alcázar, J. L., Caba, F., Zhao, C., Giancola, S., & Ghanem, B. (2022). MAD: A scalable dataset for language grounding in videos from movie audio descriptions. *Proceedings of the 2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 5026–5035. <https://doi.org/10.1109/CVPR.2014.220>
- Spong, M. W., & Vidyasagar, M. (1989). *Robot dynamics and control*. Wiley.
- Srivastava, N., Hinton, G. E., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(56), 1929–1958. <http://jmlr.org/papers/v15/srivastava14a.html>
- Srivastava, R. K., Shyam, P., Mutz, F., Jaskowski, W., & Schmidhuber, J. (2019). *Training agents using upside-down reinforcement learning*. arXiv. <https://arxiv.org/pdf/1912.02877.pdf>
- Stanić, A., Ashley, D., Serikov, O., Kirsch, L., Faccio, F., Schmidhuber, J., Hofmann, T., & Schlag, I. (2023). *The languini kitchen: Enabling language modelling research at different scales of compute*. arXiv. <https://doi.org/10.48550/arXiv.2309.11197>
- Stanić, A., Tang, Y., Ha, D., & Schmidhuber, J. (2023). Learning to generalize with object-centric agents in the open world survival game crafter. *IEEE Transactions on Evolutionary Computation*, 1–20. <https://doi.org/10.1109/TG.2023.3276849>
- Štrupl, M., Faccio, F., Ashley, D. R., Schmidhuber, J., & Srivastava, R. K. (2022). *Upside-down reinforcement learning can diverge in stochastic environments with episodic resets*. arXiv. <https://doi.org/10.48550/arXiv.2205.06595>

- Štrupl, M., Faccio, F., Ashley, D. R., Srivastava, R. K., & Schmidhuber, J. (2021). *Reward-weighted regression converges to a global optimum*. arXiv. <https://doi.org/10.48550/arXiv.2107.09088>
- Štrupl, M., Faccio, F., Ashley, D. R., Srivastava, R. K., & Schmidhuber, J. (2022). Reward-weighted regression converges to a global optimum. *Proceedings of the Thirty-Sixth AAAI Conference on Artificial Intelligence*, 8361–8369. <https://doi.org/10.1609/aaai.v36i8.20811>
- Štrupl, M., Szeher, O., Faccio, F., Ashley, D. R., Srivastava, R. K., & Schmidhuber, J. (2025). *On the convergence and stability of upside-down reinforcement learning, goal-conditioned supervised learning, and online decision transformers*. arXiv. <https://doi.org/10.48550/arXiv.2502.05672>
- Stuart, H., Wang, S., Khatib, O., & Cutkosky, M. R. (2017). The ocean one hands: An adaptive design for robust marine manipulation. *The International Journal of Robotics Research*, 36(2), 150–166. <https://doi.org/10.1177/0278364917694723>
- Sutton, R. S. (1988). Learning to predict by the methods of temporal differences. *Machine Learning*, 3, 9–44. <https://doi.org/10.1007/BF00115009>
- Sutton, R. S. (1991). Dyna, an integrated architecture for learning, planning, and reacting. *ACM SIGART Bulletin*, 2(4), 160–163. <https://doi.org/10.1145/122344.122377>
Introduced the Dyna architecture for planning in reinforcement learning.
- Sutton, R. S. (1995). Generalization in reinforcement learning: Successful examples using sparse coarse coding. *Advances in Neural Information Processing Systems*, 8, 1038–1044. <http://papers.nips.cc/paper/1109-generalization-in-reinforcement-learning-successful-examples-using-sparse-coarse-coding.pdf>
- Sutton, R. S. (2019, March 13). *The bitter lesson*. <http://incompleteideas.net/IncIdeas/BitterLesson.html>
A blog post by Sutton which argues that, in AI, more general methods always win in the long term, regardless of computational cost.
- Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning: An introduction* (1st ed.). MIT Press.
- Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction* (2nd ed.). MIT Press. <http://incompleteideas.net/book/the-book.html>
- Sutton, R. S. (1984). *Temporal credit assignment in reinforcement learning* [Doctoral dissertation, University of Massachusetts, Amherst]. <http://incompleteideas.net/papers/Sutton-PhD-thesis.pdf>
Introduced the Actor-Critic architecture.
- Taigman, Y., Yang, M., Ranzato, M., & Wolf, L. (2014). DeepFace: Closing the gap to human-level performance in face verification. *Proceedings of the 2014 IEEE Conference on Computer Vision and Pattern Recognition*, 1701–1708. <https://doi.org/10.1109/CVPR.2014.220>
- Tamar, A., Levine, S., Abbeel, P., Wu, Y., & Thomas, G. (2016). Value iteration networks. *Advances in Neural Information Processing Systems*, 29, 2146–2154. https://proceedings.neurips.cc/paper_files/paper/2016/file/c21002f464c5fc5bee3b98ced83963b8-Paper.pdf
Introduced the end-to-end differentiable Value Iteration Network architecture for planning in reinforcement learning.
- Tan, K. R., Villarino, M. L., & Maderazo, C. (2019). Automatic music mood recognition using russell’s twodimensional valence-arousal space from audio and lyrical data as classified using svm and naïve bayes. *IOP Conference Series: Materials Science and Engineering*, 482, 012019. <https://doi.org/10.1088/1757-899X/482/1/012019>
- Tang, X., Liu, Y., Cai, Z., Shao, Y., Lu, J., Zhang, Y., Deng, Z., Hu, H., An, K., Huang, R., Si, S., Chen, S., Zhao, H., Chen, L., Wang, Y., Liu, T., Jiang, Z., Chang, B., Fang, Y., ... Gerstein, M.

- (2024). *ML-Bench: Evaluating large language models and agents for machine learning tasks on repository-level code*. arXiv. <https://doi.org/10.48550/arXiv.2311.09835>
- Taylor, M. E., & Stone, P. (2009). Transfer learning for reinforcement learning domains: A survey. *Journal of Machine Learning Research*, 10(56), 1633–1685. <http://jmlr.org/papers/v10/taylor09a.html>
- Taylor, S., Park, K., Yamsani, S., & Kim, J. (2024). Fully 3d printable robot hand and soft tactile sensor based on air-pressure and capacitive proximity sensing. *Proceedings of the IEEE International Conference on Robotics and Automation*, 18100–18105. <https://doi.org/10.1109/ICRA57147.2024.10610731>
- Tesauro, G. (1994). Td-gammon, a self-teaching backgammon program, achieves master-level play. *Neural Computation*, 6(2), 215–219. <https://doi.org/10.1162/NECO.1994.6.2.215>
- Tessler, C., Givony, S., Zahavy, T., Mankowitz, D. J., & Mannor, S. (2017). A deep hierarchical approach to lifelong learning in minecraft. *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence*, 1553–1561. <https://aaai.org/ocs/index.php/AAAI/AAAI17/paper/view/14630>
- Tian, M., Gao, L., Zhang, S. D., Chen, X., Fan, C., Guo, X., Haas, R., Ji, P., Krongchon, K., Li, Y., Liu, S., Luo, D., Ma, Y., Tong, H., Trinh, K., Tian, C., Wang, Z., Wu, B., Yin, S., ... Peng, H. (2024). SciCode: A research coding benchmark curated by scientists. *Advances in Neural Information Processing Systems*, 38, 30624–30650. https://proceedings.neurips.cc/paper_files/paper/2024/file/36850592258c8c41cecdaa3dea5ff7de-Paper-Datasets_and_Benchmarks_Track.pdf
Introduced the Med QA dataset.
- Todorov, E., Erez, T., & Tassa, Y. (2012). MuJoCo: A physics engine for model-based control. *Proceedings of the 2012 IEEE/RSJ International Conference on Intelligent Robots and Systems*, 5026–5033. <https://doi.org/10.1109/IROS.2012.6386109>
- Togelius, J., Karakovskiy, S., Koutník, J., & Schmidhuber, J. (2009). Super mario evolution. *Proceedings of the 2009 IEEE Symposium on Computational Intelligence and Games*, 156–161. <https://doi.org/10.1109/CIG.2009.5286481>
- Togelius, J., & Schmidhuber, J. (2008). Computational intelligence and game design. *Proceedings of the PPSN Workshop on Computational Intelligence and Games*. <https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=658789f59e7f43f779d7fa4317f9a52a847cc73d>
- Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., Bashlykov, N., Batra, S., Bhargava, P., Bhosale, S., Bikel, D., Blecher, L., Canton-Ferrer, C., Chen, M., Cucurull, G., Esiobu, D., Fernandes, J., Fu, J., Fu, W., ... Scialom, T. (2023). *Llama 2: Open foundation and fine-tuned chat models*. arXiv. <https://doi.org/10.48550/arXiv.2307.09288>
Introduced the Llama 2 family of LLMs.
- Turing, A. M. (1950). Computing machinery and intelligence. *Mind*, 59(236), 433–460. <https://doi.org/10.1093/mind/LIX.236.433>
- United Nations, Department of Economic and Social Affairs, Population Division. (2019). *World population prospects 2019: Highlights*. https://population.un.org/wpp/Publications/Files/WPP2019_Highlights.pdf
- Utter, A., Zakrzewski, M., Keene, A., Dietrich, S., Lin, S., McDonald, E., Whitehair, N., & Zheng, J. (2020). SatCat5: A low-power, mixed-media ethernet network for smallsats. *Proceedings of the 34th Annual Small Satellite Conference*. <http://digitalcommons.usu.edu/smallsat/2020/all2020/174/>
- Vall, A., Quadrana, M., Schedl, M., & Widmer, G. (2019). Order, context and popularity bias in next-song recommendations. *International Journal of Multimedia Information Retrieval*, 8(2), 101–113. <https://doi.org/10.1007/s13735-019-00169-8>

- VandenBos, G. R. (2015). *Apa dictionary of psychology* (2nd ed.). American Psychological Association. <https://doi.org/10.1037/14646-000>
- Vasan, G. (2017). *Teaching a Powered Prosthetic Arm with an Intact Arm Using Reinforcement Learning* [Master's thesis, University of Alberta]. <https://doi.org/10.7939/R3GFoN94H>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998–6008. <http://papers.nips.cc/paper/7181-attention-is-all-you-need.pdf>
Introduced the transformer architecture.
- Ventura, J. D., Ashley, D. R., Herrmann, V., Faccio, F., & Schmidhuber, J. (2025). *Upside down reinforcement learning with policy generators*. arXiv. <https://doi.org/10.48550/arXiv.2501.16288>
- Vinyals, O., Babuschkin, I., Czarnecki, W. M., Mathieu, M., Dudzik, A., Chung, J., Choi, D. H., Powell, R., Ewalds, T., Georgiev, P., Oh, J., Horgan, D., Kroiss, M., Danihelka, I., Huang, A., Sifre, L., Cai, T., Agapiou, J. P., Jaderberg, M., ... Silver, D. (2019). Grandmaster level in starcraft II using multi-agent reinforcement learning. *Nature*, 575(7782), 350–354. <https://doi.org/10.1038/S41586-019-1724-Z>
- Vinyals, O., Ewalds, T., Bartunov, S., Georgiev, P., Vezhnevets, A. S., Yeo, M., Makhzani, A., Küttler, H., Agapiou, J. P., Schrittwieser, J., Quan, J., Gaffney, S., Petersen, S., Simonyan, K., Schaul, T., Hasselt, H. van, Silver, D., Lillicrap, T. P., Calderone, K., ... Tsing, R. (2017). *Starcraft II: A new challenge for reinforcement learning*. arXiv. <https://doi.org/10.48550/arXiv.1708.04782>
- Vonnegut, K. (1981). *Palm sunday*. Delacorte Press.
- Wang, J., Hu, X., Hou, W., Chen, H., Zheng, R., Wang, Y., Yang, L., Huang, H., Ye, W., Geng, X., Jiao, B., Zhang, Y., & Xie, X. (2023). *On the robustness of ChatGPT: An adversarial and out-of-distribution perspective*. arXiv. <https://doi.org/10.48550/arXiv.2302.12095>
- Wang, P., Li, L., Chen, L., Cai, Z., Zhu, D., Lin, B., Cao, Y., Kong, L., Liu, Q., Liu, T., & Sui, Z. (2024). Large language models are not fair evaluators. *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 9440–9450. <https://doi.org/10.18653/v1/2024.acl-long.511>
- Wang, T., Kulikov, I., Golovneva, O., Yu, P., Yuan, W., Dwivedi-Yu, J., Pang, R. Y., Fazel-Zarandi, M., Weston, J., & Li, X. (2024). *Self-taught evaluators*. arXiv. <https://doi.org/10.48550/arXiv.2408.02666>
- Wang, W., Alyahya, H. A., Ashley, D. R., Serikov, O., Khizbullin, D., Faccio, F., & Schmidhuber, J. (2024). *How to correctly do semantic backpropagation on language-based agentic systems*. arXiv. <https://doi.org/10.48550/arXiv.2412.03624>
- Wang, Y., Ma, X., Zhang, G., Ni, Y., Chandra, A., Guo, S., Ren, W., Arulraj, A., He, X., Jiang, Z., Li, T., Ku, M., Wang, K., Zhuang, A., Fan, R., Yue, X., & Chen, W. (2024). MMLU-Pro: A more robust and challenging multi-task language understanding benchmark. *Advances in Neural Information Processing Systems*, 37, 95266–95290. https://papers.nips.cc/paper_files/paper/2024/file/ad236edc564f3e3156e1b2feafb99a24-Paper-Datasets_and_Benchmarks_Track.pdf
- Wang, Y., Wu, Q., Ashley, D. R., Faccio, F., Li, W., Huang, C., & Schmidhuber, J. (2025). Scaling value iteration networks to 5000 layers for extreme long-term planning. *Proceedings of the 42nd International Conference on Machine Learning*, 267, 64914–64936. <https://raw.githubusercontent.com/mlresearch/v267/main/assets/wang25do/wang25do.pdf>
- Wei, J., Bosma, M., Zhao, V. Y., Guu, K., Yu, A. W., Lester, B., Du, N., Dai, A. M., & Le, Q. V. (2022). Finetuned language models are zero-shot learners. *Proceedings of the 10th International Conference on Learning Representations*. <https://openreview.net/pdf?id=gEZrGCozdqR>

- Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., Chi, E. H., Hashimoto, T., Vinyals, O., Liang, P., Dean, J., & Fedus, W. (2022). Emergent abilities of large language models. *Transactions on Machine Learning Research*. <https://openreview.net/pdf?id=yzkSU5zdwD>
- Wei, J., Yang, C., Song, X., Lu, Y., Hu, N., Huang, J., Tran, D., Peng, D., Liu, R., Huang, D., Du, C., & Le, Q. V. (2024). Long-form factuality in large language models. *Advances in Neural Information Processing Systems*, 37, 80756–80827. https://proceedings.neurips.cc/paper_files/paper/2024/file/937ae0e83e808d2cb8627feidef8c751-Paper-Conference.pdf
Introduced the Med QA dataset.
- Wierstra, D., Förster, A., Peters, J., & Schmidhuber, J. (2007). Solving deep memory pomdps with recurrent policy gradients. *Proceedings of the 17th International Conference on Artificial Neural Networks*, 4668, 697–706. https://doi.org/10.1007/978-3-540-74690-4_71
- Wierstra, D., Förster, A., Peters, J., & Schmidhuber, J. (2010). Recurrent policy gradients. *Logic Journal of the IGPL*, 5(18), 620–634. <https://doi.org/10.1093/JIGPAL/JZPo49>
- Wierstra, D., Schaul, T., Peters, J., & Schmidhuber, J. (2008a). Episodic reinforcement learning by logistic reward-weighted regression. *Proceedings of the 18th International Conference on Artificial Neural Networks*, 5163, 407–416. https://doi.org/10.1007/978-3-540-87536-9_42
- Wierstra, D., Schaul, T., Peters, J., & Schmidhuber, J. (2008b). Fitness expectation maximization. *Proceedings of the 10th International Conference on Parallel Problem Solving from Nature*, 5199, 337–346. https://doi.org/10.1007/978-3-540-87700-4_34
- Witten, I. H. (1977). An adaptive optimal controller for discrete-time markov environments. *Information and control*, 34(4), 286–295. [https://doi.org/10.1016/S0019-9958\(77\)90354-0](https://doi.org/10.1016/S0019-9958(77)90354-0)
- Wolpert, D. M., Ghahramani, Z., & Flanagan, J. R. (2001). Perspectives and problems in motor learning. *Trends in Cognitive Sciences*, 5(11), 487–494. [https://doi.org/10.1016/S1364-6613\(00\)01773-3](https://doi.org/10.1016/S1364-6613(00)01773-3)
- Wu, C., Yin, S., Qi, W., Wang, X., Tang, Z., & Duan, N. (2023). *Visual ChatGPT: Talking, drawing and editing with visual foundation models*. arXiv. <https://doi.org/10.48550/arXiv.2303.04671>
- Wu, Q., Bansal, G., Zhang, J., Wu, Y., Zhang, S., Zhu, E., Li, B., Jiang, L., Zhang, X., & Wang, C. (2024). *AutoGen: Enabling next-gen LLM applications via multi-agent conversation framework*. arXiv. <https://doi.org/10.48550/arXiv.2308.08155>
- Xiao, H., Rasul, K., & Vollgraf, R. (2017). *Fashion-MNIST: A novel image dataset for benchmarking machine learning algorithms*. arXiv. <https://doi.org/10.48550/arXiv.1708.07747>
- Yu, Y., Yao, Z., Li, H., Deng, Z., Jiang, Y., Cao, Y., Chen, Z., Suchow, J. W., Cui, Z., Liu, R., Xu, Z., Zhang, D., Subbalakshmi, K., Xiong, G., He, Y., Huang, J., Li, D., & Xie, Q. (2024). Fincon: A synthesized LLM multi-agent system with conceptual verbal reinforcement for enhanced financial decision making. *Advances in Neural Information Processing Systems*, 37, 137010–137045. https://proceedings.neurips.cc/paper_files/paper/2024/file/f7ae4fe91d96f5oabc2211fo9b6a7e49-Paper-Conference.pdf
- Zeng, Y., Lin, H., Zhang, J., Yang, D., Jia, R., & Shi, W. (2024). *How johnny can persuade LLMs to jailbreak them: Rethinking persuasion to challenge AI safety by humanizing LLMs*. arXiv. <https://doi.org/10.48550/arXiv.2401.06373>
Shows how persuasion techniques from the social sciences can be used to jailbreak LLMs effectively.
- Zenke, F., Poole, B., & Ganguli, S. (2017). Continual learning through synaptic intelligence. *Proceedings of the 34th International Conference on Machine Learning*, 70, 3987–3995. <http://proceedings.mlr.press/v70/zenker17a/zenker17a.pdf>

- Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2017). Understanding deep learning requires rethinking generalization. *Proceedings of the 5th International Conference on Learning Representations*. <https://openreview.net/forum?id=Sy8gdB9xx>
- Zhao, C., Chang, E., Liu, Z., Chang, C., Wen, W., Lai, C., Cao, S., Tian, Y., Krishnamoorthi, R., Shi, Y., & Chandra, V. (2025). *MobileLLM-R1: Exploring the limits of sub-billion language model reasoners with open training recipes*. arXiv. <https://doi.org/10.48550/arXiv.2509.24945>
- Zheng, L., Chiang, W., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., & Stoica, I. (2023). Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36, 46595–46623. https://proceedings.neurips.cc/paper_files/paper/2023/file/91f18a1287b398d378ef22505bf41832-Paper-Datasets_and_Benchmarks.pdf
- Zheng, Q., Zhang, A., & Grover, A. (2022). *Online decision transformer*. arXiv. <https://doi.org/10.48550/arXiv.2202.05607>
- Zhu, H., & Feron, E. (2021). A systems engineering approach to the design and education of a robotic baby. *Proceedings of the 31st INCOSE International Symposium*, 31(1), 327–342. <https://incose.onlinelibrary.wiley.com/doi/pdf/10.1002/j.2334-5837.2021.00841.x>
- Zhu, H., Wilson, S., & Feron, E. (2023). The design, education and evolution of a robotic baby. *IEEE Transactions on Robotics*, 39(3), 2488–2507. <https://doi.org/10.1109/TRO.2023.3240619>
- Zhuge, M., Liu, H., Faccio, F., Ashley, D. R., Csordás, R., Gopalakrishnan, A., Hamdi, A., Hammoud, H. A. A. K., Herrmann, V., Irie, K., Kirsch, L., Li, B., Li, G., Liu, S., Mai, J., Piekos, P., Ramesh, A., Schlag, I., Shi, W., ... Schmidhuber, J. (2023). *Mindstorms in natural language-based societies of mind*. arXiv. <https://doi.org/10.48550/arXiv.2305.17066>
- Zhuge, M., Liu, H., Faccio, F., Ashley, D. R., Csordás, R., Gopalakrishnan, A., Hamdi, A., Hammoud, H. A. A. K., Herrmann, V., Irie, K., Kirsch, L., Li, B., Li, G., Liu, S., Mai, J., Piekos, P., Ramesh, A., Schlag, I., Shi, W., ... Schmidhuber, J. (2025). Mindstorms in natural language-based societies of mind. *Computational Visual Media*, 11(1), 29–81. <https://doi.org/10.26599/CVM.2025.9450460>
- Zhuge, M., Zhao, C., Ashley, D. R., Wang, W., Khizbullin, D., Xiong, Y., Liu, Z., Chang, E., Krishnamoorthi, R., Tian, Y., Shi, Y., Chandra, V., & Schmidhuber, J. (2025). Agent-as-a-Judge: Evaluate agents with agents. *Proceedings of the 42nd International Conference on Machine Learning*, 267, 80569–80611. <https://raw.githubusercontent.com/mlresearch/v267/main/assets/zhuge25a/zhuge25a.pdf>
- Zou, A., Wang, Z., Kolter, J. Z., & Fredrikson, M. (2023). *Universal and transferable adversarial attacks on aligned language models*. arXiv. <https://doi.org/10.48550/arXiv.2307.15043>
Proposed the random search-based Greedy Coordinate Gradient (GCG) algorithm to produce a single adversarial suffix that jailbreaks LLMs.
- Zuckerberg, M. E. (2025). *Personal superintelligence for everyone*. Meta Platforms, Inc. <https://www.meta.com/superintelligence/>
Simple text post where Zuckerberg says superintelligence is in sight.